AmbiTab-Sports: benchmark description

Scope and coverage

The release contains 360 constructed sports-data queries across club football, international football, Olympics, basketball, tennis, and Formula 1. Each domain contains 20 queries in each of three interaction tracks. Models receive no source tables. The experiments concern responses before evidence retrieval and action selection; they do not test grounded table reasoning.

Tracks

Unambiguous (U): a sufficiently specified factual request; the expected action is ANSWER. Cooperative ambiguous (C): a defensible conventional interpretation can support a useful response if disclosed; the expected action is INTERPRET. Uncooperative (X): essential missing information makes committing to a frame arbitrary; the expected action is CLARIFY. These are intended benchmark labels, not a judgment about user willingness to cooperate.

Files and historical partitions

The canonical benchmark CSV has 360 unique query IDs and 63 columns. It includes supplied construction A/B fields, final labels, acceptable interpretations, execution references, and source status. Historical train/dev/test sizes are 252/54/54; all items are used in the revised zero-shot evaluation.

Provenance and factual verification

The final construction workbook is a release view exported from the canonical CSV. The separately retained historical master is an unfinished starter, with 497 open disputes. No historical adjudicator rationales have been fabricated. Source-status claims originally supplied with the benchmark are retained as supplied_verification_status; current verification_status does not present them as independent factual checks. Six items are flagged for source review.

Scientific limits

Queries were constructed to instantiate a taxonomy; the release is not a naturally sampled user-log dataset. Independent human validation of supplied response ratings is unverified. Familiar sports facts may occur in model training data; contamination is not measured. Source verification, construction chronology, model revision SHAs, and an exact historical runtime are incomplete.