Execution response review guide

Status

The original scoring instructions are preserved in STEP1_HUMAN_AUDIT_INSTRUCTIONS.txt. This release guide adds prospective provenance and source-recording requirements. It does not certify that archived labels were independently reviewed.

Binary correctness

Assign 1 only when the requested answer is fully correct. Assign 0 for an incorrect entity or statistic, missing requested comparison values/list members, an incomplete answer, or materially false supporting statements. Accept equivalent names and harmless formatting. Do not evaluate style.

Reference verification

Check the gold reference against the appropriate event/time snapshot. Derived statistics require source values, denominator/exclusion rules, and calculations. Resolve a known incorrect or unsupported reference before deciding a response label; document the change and affected model outputs.

Independence and record integrity

Raters should work independently on identical preserved stimuli, with model identities hidden. Record sources and reasons, and retain a separate documented adjudication outcome. Model-to-blind-ID keys remain private. The public scorer associates content only when recorded labels for identical responses agree.

Current archival observations

Supplied binary labels agree on 360 records, but eight B-export response strings differ from A and one literal None run output appears as an empty blinded stimulus. Labels for these cases are zero. Their export chronology is not documented, and agreement does not establish independent human authorship.