Saved-response evaluation protocol

Reproducible analysis

Run python code/reproduce.py from the repository. It reads included source files and writes regenerated results/figures under reproduced/. No network, GPU, model download, or new annotation is required for saved-response analysis. Historical inference is archived separately and not rerun by this command.

Execution

Analyze 120 U queries for each of three LLMs, using supplied binary correctness records. Equivalent names and harmless formatting are accepted by the supplied instructions; incomplete requested lists/comparisons and materially false supporting facts are incorrect. Model identities are associated by query and response content, without publishing private blind-ID keys. One literal None response is matched to its empty blinded export; the original text and normalization flag remain in the scored output.

Cooperative quality

The supplied final corpus contains 480 responses, four systems on 120 C queries. Three recorded dimensions range from 0 to 2; total/6 is the normalized score. Agreement-source rows number 205; 275 are marked adjudicated. These labels do not establish independent human review or independently verified factual correctness. The proposed prospective re-review rubric is separate from the unverified historical protocol.

Unified action selection

All 360 queries receive the same action prompt, with track labels hidden. Score ANSWER, INTERPRET, or CLARIFY against the intended action. Report accuracy, balanced accuracy, macro-F1, per-track recall, appropriate X clarification, U/C over-clarification, and confusion matrices. X non-clarification is a policy mismatch; it is not a measured hallucination or harm rate. Always-clarify accuracy is 1/3, with macro-F1 1/6.

Inference and uncertainty

The recorded repositories are Qwen/Qwen2.5-7B-Instruct, google/gemma-4-E4B-it, and mistralai/Mistral-7B-Instruct-v0.2. Historical commit SHAs are not recorded. Archived settings describe greedy generation, an input budget of 2048 tokens, and output budgets of 100/140/80 for execution/action/legacy conditions. These do not establish a fully locked historical environment or explicit thinking-mode default.

Statistics

Bootstrap means and paired differences use 5000 resamples at seed 42. Archived CSV row order is retained. Execution and action accuracy use exact two-sided McNemar tests; cooperative scores use 10000 paired sign-flip draws. Holm correction covers the three LLM pairs separately within each primary family. Non-significance is not equivalence; subgroup and dimension comparisons remain descriptive.

Primary versus secondary evidence

The corrected master uses recorded binary execution labels, cooperative rubric scores, and unified action metrics. Task-aware automatic execution scoring and legacy clarification/abstention weighting are secondary. The original rank-reversal claim is withdrawn. Old supplied summaries are preserved only under archive/supplied_results.