Cooperative response rubric for prospective validation

Status and purpose

This is a proposed operational guide for a future independent human re-review. It has not been applied to the archived scores and is not claimed to be their historical annotation protocol. A validated replacement for the rejected automatic judge still requires actual independent review. Use anonymous rater identifiers and keep model identities hidden while rating.

Work independently

Each rater reviews every assigned response without access to another rating, a model name, or a suggested machine score. Record the interpretation, assumptions, sources checked, uncertainty, and rating reasons. Disagreements go to a separately documented adjudication step; preserve both original ratings and the final reason.

Interpretation validity: 0 to 2

0: the chosen frame is unrelated, arbitrary, or fails to resolve the material ambiguity. 1: the frame is plausible but an important part of the query remains unresolved or the default is weakly justified. 2: the response chooses a coherent, conventional frame that makes the request answerable and explains why it is defensible.

Assumption quality: 0 to 2

0: material assumptions are hidden, contradictory, or unsupported. 1: assumptions are partly stated but a relevant entity, time window, metric, or scope remains unclear. 2: the response explicitly states the assumptions needed for its frame, bounds the time/competition/statistic, and communicates relevant uncertainty or alternatives.

Answer usefulness: 0 to 2

0: no substantive answer is given, the answer does not address the chosen frame, or a major factual error defeats the request. 1: the answer is relevant and partly useful but incomplete, insufficiently supported, or missing a requested component. 2: the answer directly and completely addresses the chosen frame and material factual claims have been checked. Do not award 2 solely because the answer is fluent or contains specific numbers.

Factual checks and equivalent wording

Accept semantically equivalent names and harmless formatting. Use primary sports statistics where available; preserve source URL, event/time snapshot, relevant data, and calculation for derived statistics. A known source mismatch must remain unresolved rather than being marked verified. Missing tables in the experimental condition should not be confused with permission to invent facts.

Anchors

A response that only says “assume the latest season and use wins” without an answer can receive partial framing credit but no complete usefulness credit. A response that identifies a specific season/competition, discloses a conventional statistic, and gives a checked answer can receive full credit. A response giving wrong championship totals cannot receive full factual-usefulness credit even when its format is correct.

Data recording

Preserve the original response text exactly, including signs, percentages, and score strings. Avoid spreadsheet conversion of results such as 3-1 into dates. Record three integers, reasons, sources, and adjudication status separately. Preserve blank/missing output distinctions.

Validation after review

Report exact agreement and quadratic Cohen kappa separately per dimension before adjudication. Compare the proposed judge to an independently rated validation subset before using it as a primary model-ranking metric. Recompute all primary results if labels change, version the updated corpus, and report disagreements rather than overwriting history.